Brain Tumor Segmentation with Random Forest and U-Net

Shuhan Xiao and Alexander Kugele

Heidelberg University

Motivation

- Brain tumors need immediate treatment
- Even experts can not segment perfectly
- Takes time to go through a complete MR scan

Efficiently segmenting tumors automatically improves treatment planing

Dataset

- 3D MR scans of 275 human brains
- 4 scan types, T1, T1c, T2 and Flair
- depth = 155, height = 240, width = 240
- ⇒ 170500 images
- 5 classes (we only use 2)
- center-cropped to 80% of their size

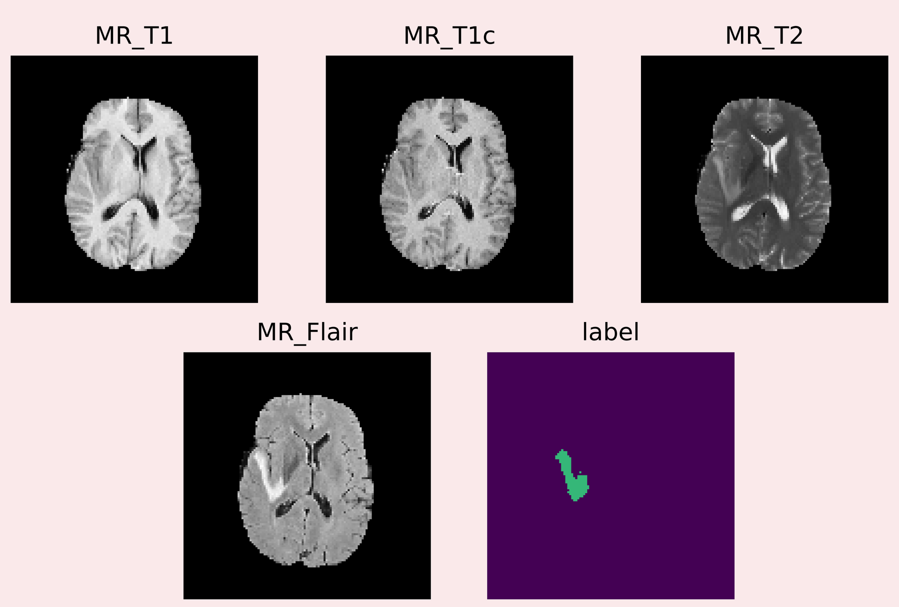

Features

- Gaussian, LoG, Gaussian gradient
- Hessian and structure tensor eigenvalues
- equalized histogram
- 29 features in total
- 1k images + features ≈ 30 GB of memory

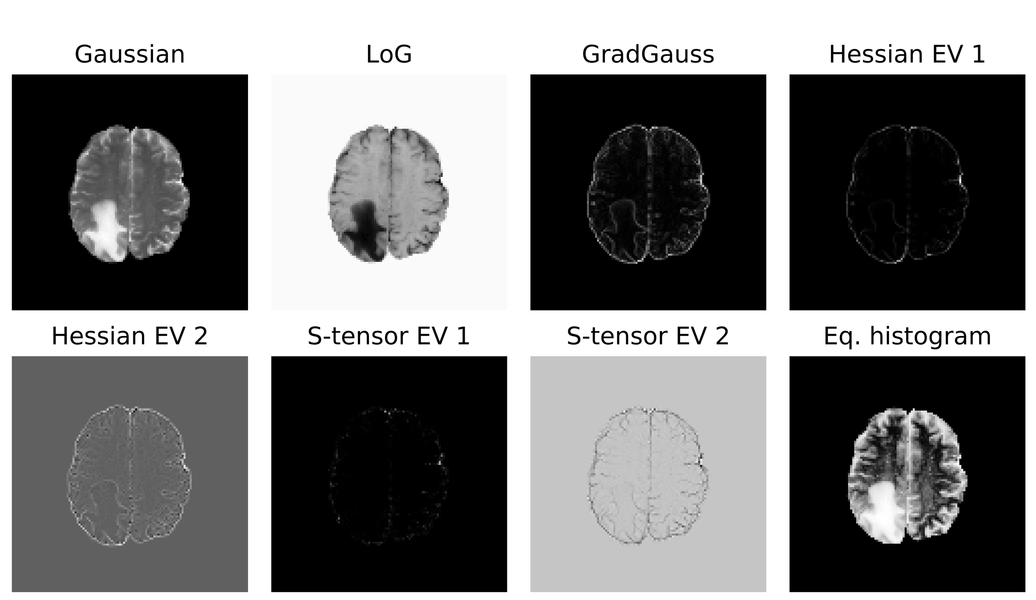

Random Forest [1]

- default settings of scikit-learn implementation
- trained in batch-mode
- 100 estimators per batch, 3 batches of 1k images

Advantages and disadvantages

- ➕ Simple
- ➕ Rarely overfits
- ➕ Variable image input size
- ➕ Optimized implementation available
- ➖ No GPU training
- ➖ No incremental training (for vanilla RF)
- ➖ Features have to be hand-selected

U-Net [2]

modifications: depth = 4, loss = 1 - Dice, padding

Training: 10 images/batch, 30 epochs

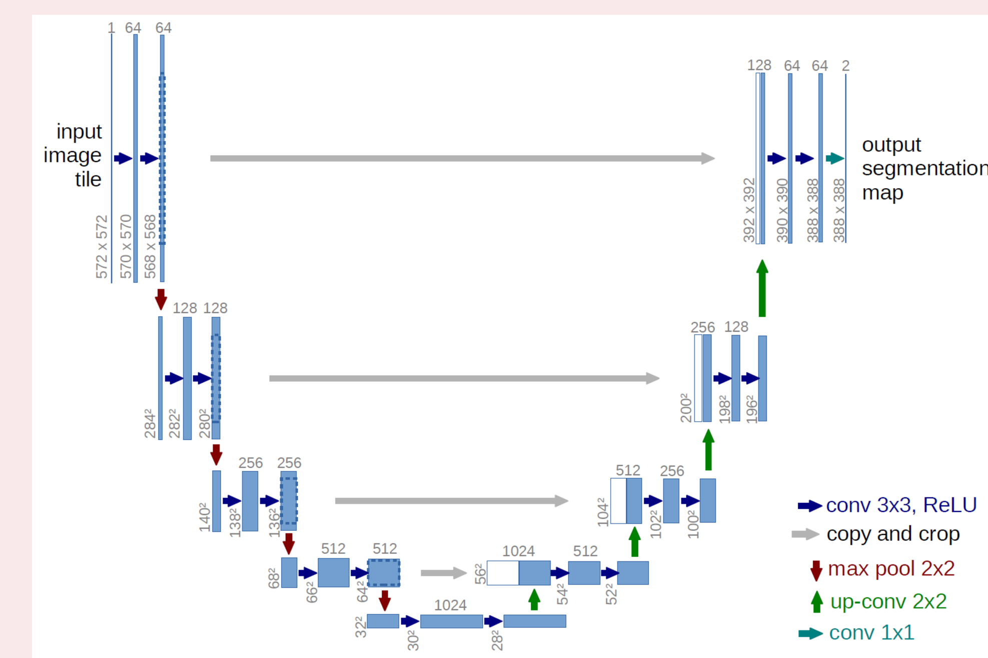

Final Scores

	Dice [%]	Sensitivity [%]	Specificity [%]
Random Forest	65.9	77.1	99.1
U-Net	69.3	73.1	

RF: Results

- training time ≈ 13 h
- inference time ≈ 2 − 20s per image
- load time ≈ 310s
- disk space ≈ 15 GB (pickled)
- memory space during inference ≈ 20 GB

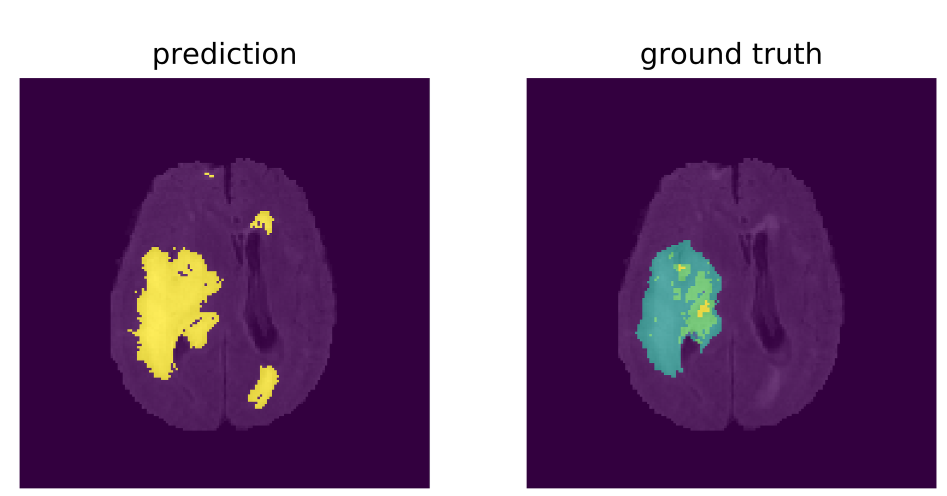

U-Net: Results

- training time ≈ 1.5 h
- inference time ≈ 0.5s per image
- load time ≈ 20s
- disk space ≈ 0.3 GB
- memory space during inference ≈ 0.7 GB

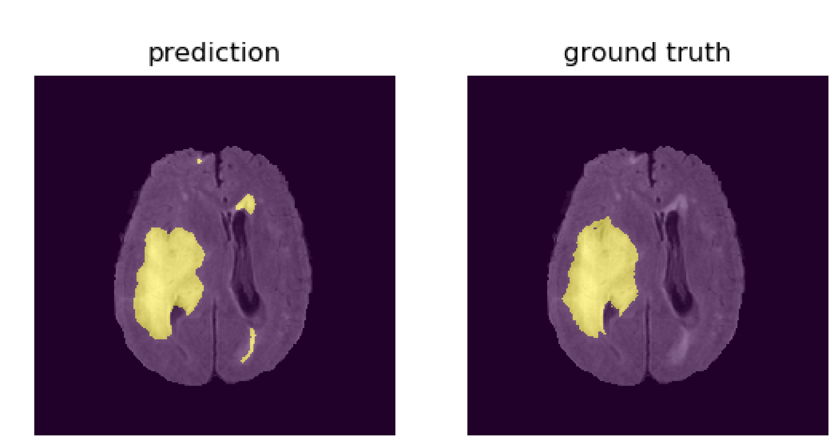

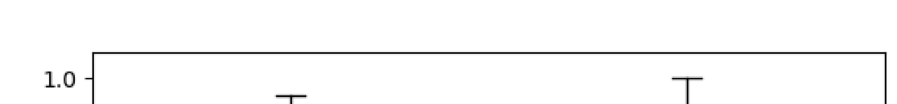

Loss

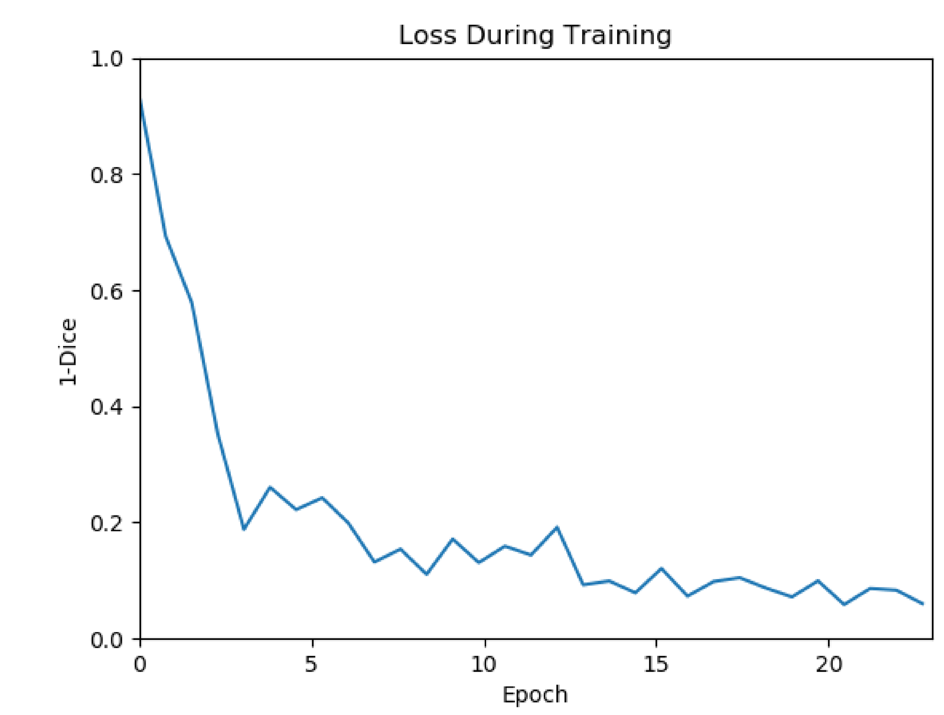

Conclusion

- Both methods achieve similar results regarding the scores
- Random Forest is easier to set up
- U-Net is easier and faster to train
- None of the two methods could achieve state-of-the-art results out of the box

References

[1] M. Ristin, M. Guillaumin, J. Gall, and L. Van Gool. Incremental learning of random forests for large-scale image classification. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 38(3):490–503, March 2016.

[2] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *International Conference on Medical image computing and computer-assisted intervention*, pages 234–241. Springer, 2015.

UNIVERSITÄT HEIDELBERG
ZUKUNFT SEIT 1386